

# Homework 5

Stat 428

Homework submission guideline: For the convenience of grading, you are required to submit the **Rmd** and **PDF** files separately and NOT zipped, as a zip file cannot be previewed in Canvas. You may get a penalty if wrong format is submitted.

Before the homework, please set seed as your UIN for the sake of reproducibility.

```
set.seed(677978239)
```

## 1. Cross Validation (40 points)

We're going to look at a data set on 97 men who have prostate cancer (from the book The Elements of Statistical Learning). There are 10 variables measured on these 97 men:

1. **lpsa**: log PSA score
2. **lcavol**: log cancer volume
3. **lweight**: log prostate cancer weight
4. **age**: age of patient
5. **lbph**: log of the amount of benign prostatic hyperplasia
6. **svi**: seminal vesicle invasion
7. **lcp**: log of capsular penetration
8. **gleason**: Gleason score
9. **pgg45**: percent of Gleason scores 4 or 5
10. **train**: if belonging to training data set

To load this prostate cancer data set and store it as a matrix **pros.data**, we can do as following:

```
pros.data =  
  as.matrix(read.table("https://web.stanford.edu/~hastie/ElemStatLearn/datasets/prostate.data"))
```

We only use the first 9 columns

```
pros.data = pros.data[,1:9]
```

In this question, we are going to predict **lpsa** by a linear model of other variables. In particular, we consider the following three models:

$$X_{lpsa} = \beta_0 + \beta_1 X_{lcavol} + \beta_2 X_{lweight} + \beta_3 X_{age} + \epsilon$$

$$X_{lpsa} = \beta_0 + \beta_1 X_{lcavol} + \beta_2 X_{lweight} + \epsilon$$

$$X_{lpsa} = \beta_0 + \beta_1 X_{lcavol} + \epsilon$$

- **1a.** Randomly split the prostate cancer data frame into  $k = 5$  folds of roughly equal size. (Make sure your code is general enough to handle an arbitrary number of folds  $k$ ; you will be asked to change the number of folds in questions that follow.) Report the number of observations that fall in each fold.

```
summary(pros.data)
```

```
##      lcavol      lweight      age      lbph
## Min.   :-1.3471  Min.    :2.375  Min.    :41.00  Min.    :-1.3863
## 1st Qu.: 0.5128  1st Qu.:3.376  1st Qu.:60.00  1st Qu.: -1.3863
## Median : 1.4469  Median :3.623  Median :65.00  Median : 0.3001
## Mean   : 1.3500  Mean    :3.629  Mean    :63.87  Mean    : 0.1004
## 3rd Qu.: 2.1270  3rd Qu.:3.876  3rd Qu.:68.00  3rd Qu.: 1.5581
## Max.    : 3.8210  Max.    :4.780  Max.    :79.00  Max.    : 2.3263
##      svi      lcp      gleason      pgg45
## Min.    :0.0000  Min.    :-1.3863  Min.    :6.000  Min.    : 0.00
## 1st Qu.:0.0000  1st Qu.: -1.3863  1st Qu.:6.000  1st Qu.: 0.00
## Median :0.0000  Median :-0.7985  Median :7.000  Median : 15.00
## Mean    :0.2165  Mean    :-0.1794  Mean    :6.753  Mean    : 24.38
## 3rd Qu.:0.0000  3rd Qu.: 1.1787  3rd Qu.:7.000  3rd Qu.: 40.00
## Max.    :1.0000  Max.     :2.9042  Max.     :9.000  Max.     :100.00
##      lpsa
## Min.    :-0.4308
## 1st Qu.: 1.7317
## Median : 2.5915
## Mean    : 2.4784
## 3rd Qu.: 3.0564
## Max.    : 5.5829
```

```
nfolds = 5
folds <- split(sample(nrow(pros.data), nrow(pros.data), replace = FALSE), as.factor(1:nfolds))
datalist <- lapply(folds, function(x) pros.data[x, ])
i = 0
n = nrow(pros.data)

for(data in datalist) {
  print(nrow(data))
}
```

```
## [1] 20
## [1] 20
## [1] 19
## [1] 19
## [1] 19
```

As we can see, the length of the 5 folds are roughly equal.

- **1b.** Over the folds you computed in the previous question, compute the cross-validation error of the above three linear models.

```

e1 <- e2 <- e3 <- numeric(nfolds)
for(i in 1:nfolds){
  A = datalist[[i]]

  B = data.frame()
  for (j in 1:nfolds) {
    if (i != j) {
      B = rbind(B, datalist[[j]])
    }
  }
  A.lpsa <- A[, "lpsa"]
  A.lcavol = A[, "lcavol"]
  A.lweight = A[, "lweight"]
  A.age = A[, "age"]

  A = as.data.frame(A)
  B = as.data.frame(B)

  L1 <- lm(lpsa ~ lcavol+lweight+age, data=B);
  hatAlpsa=L1$coef[1]+L1$coef[2]*A.lcavol + L1$coef[3]*A.lweight + L1$coef[4]*A.age
  e1[i]=sum((hatAlpsa-A.lpsa)^2)

  #2
  L2 <- lm(lpsa ~ lcavol+lweight, data=B);
  hatAlpsa=L2$coef[1]+L2$coef[2]*A.lcavol + L2$coef[3]*A.lweight
  e2[i]=sum((hatAlpsa-A.lpsa)^2)

  #3
  L3 <- lm(lpsa ~ lcavol, data=B);
  hatAlpsa=L3$coef[1]+L3$coef[2]*A.lcavol
  e3[i]=sum((hatAlpsa-A.lpsa)^2)
}
c(sum(e1)/n,sum(e2)/n,sum(e3)/n)

```

```
## [1] 0.5676379 0.5622533 0.6224067
```

- **1c.** Write a function `pros.cv()`, which takes three arguments: `df`, a data frame of prostate cancer measurements, with a default of `pros.df`; `k`, an integer determining the number of cross-validation folds, with a default of 5; and `seed`, an integer to be passed to `set.seed()` before defining the folds, with a default of `NULL` (meaning no seed shall be set). Your function should split up the given data `df` into `k` folds of roughly equal size, and using these folds, compute the cross-validation error of the above three linear model. Its output should simply be a vector of cross-validation errors. When `k` is equal to the row of the data set, leave-one-out cross validation should be used.

```

pros.cv <- function(df, k=5, seed=NULL) {
  if (!is.null(seed)) {
    set.seed(seed)
  }

  nfolds = k
  folds <- split(sample(nrow(df), nrow(df), replace = FALSE), as.factor(1:nfolds))
  datalist <- lapply(folds, function(x) df[x, ])
  i = 0

```

```

e1 <- e2 <- e3 <- numeric(nfolds)
n = nrow(pros.data)
for(i in 1:nfolds){
  A = data.frame(matrix(nrow=1, ncol = ncol(pros.data)))
  colnames(A) = colnames(pros.data)
  A = rbind(A, datalist[[i]])
  B = data.frame(matrix(nrow=1, ncol = ncol(A)))
  colnames(B) = colnames(A)

  for (j in 1:nfolds) {
    if (i != j) {
      B = rbind(B, datalist[[j]])
    }
  }
  A = na.omit(A)
  B = na.omit(B)

  A.lpsa <- A$lpsa
  A.lcavol = A$lcavol
  A.lweight = A$lweight
  A.age = A$age

  L1 <- lm(lpsa ~ lcavol+lweight+age, data=B);

  hatAlpsa=L1$coef[1]+L1$coef[2]*A.lcavol + L1$coef[3]*A.lweight + L1$coef[4]*A.age
  e1[i]=sum((hatAlpsa-A.lpsa)^2)

  #2
  L2 <- lm(lpsa ~ lcavol+lweight, data=B);

  hatAlpsa=L2$coef[1]+L2$coef[2]*A.lcavol + L2$coef[3]*A.lweight

  e2[i]=sum((hatAlpsa-A.lpsa)^2)

  #3

  L3 <- lm(lpsa ~ lcavol, data=B);

  hatAlpsa=L3$coef[1]+L3$coef[2]*A.lcavol

  e3[i]=sum((hatAlpsa-A.lpsa)^2)

  #Use the test and train data partitions however you desire...
}
return(c(sum(e1)/n,sum(e2)/n,sum(e3)/n))
}
pros.cv(pros.data,5)

```

```
## [1] 0.5873896 0.5868482 0.6583445
```

- **1d.** Investigate the result of `pros.cv()` for different values of `k`, specifically, for `k` equal to 2, 4, 8, 16

and 97. For each value, run `pros.cv()` some large number of times (say, 50) and report the average of the cross-validation error estimates, and the standard deviation of these estimates. Then, plot them in an informative way (say, a box plot with `boxplot()`). What do you notice? Is this surprising?

```
lst = c(2,4,8,16,97)
```

```
average = list()
```

```
std = list()
```

```
val = list()
```

```
print(nrow(pros.data))
```

```
## [1] 97
```

```
index = 1
```

```
for(k in lst) {
```

```
  first = rep(0, 50)
```

```
  second = rep(0, 50)
```

```
  third = rep(0, 50)
```

```
  for(i in 1:50) {
```

```
    temp = pros.cv(pros.data, k)
```

```
    first[i] = temp[1]
```

```
    second[i] = temp[2]
```

```
    third[i] = temp[3]
```

```
  }
```

```
  boxplot(first,second, third)
```


```
  val[index] <- c((first), (second), (third))
```


```
  average[index] <- c(mean(first), mean(second), mean(third))
```


```
  std[index] <- c(sd(first), sd(second), sd(third))
```

```
  index = index + 1
```


```
}
```














```
print("Average:")
```

```
## [1] "Average:"
```

```
print(unlist(average))
```

```
## [1] 0.6053231 0.5841892 0.5719804 0.5740478 0.5729621
```

```
print("STD")
```

```
## [1] "STD"
```

```
print(unlist(std))
```

```
## [1] 7.116361e-02 2.637741e-02 1.327123e-02 8.620154e-03 1.154649e-16
```

As the  $k$  gets bigger, the std get smaller while the average have no significant changes. It does not surprise to me

## 2. Permutation Tests (60 points)

- **2a.** The nearest neighbor test can also be applied to one dimensional data. Can you design a simulation experiment to compare Type I and Type II error of Kolmogorov-Smirnov test and nearest neighbor test on a one-dimensional data set?

Here is the ks.perm function that I got from the lecture.

```
ks.perm <- function(X,Y,alpha,B) {  
  D=ks.test(X,Y)$statistic  
  Z=c(X,Y); N=1:length(Z)  
  DPerm=rep(0,B)  
  for (b in 1:B) {  
    selectedindex <- sample(N, size = length(X), replace = FALSE)  
    X1 <- Z[selectedindex]  
    Y1 <- Z[-selectedindex]  
    DPerm[b]=ks.test(X1,Y1)$statistic  
  }  
  P <- mean(c(D, DPerm) >= D)  
  return(P<alpha)  
}
```

Check the Type I error:

```
n = 100; m=100  
times = 100  
decision <- numeric(times)  
for (t in 1:times) {  
  X=rnorm(n)  
  Y=rnorm(m)  
  decision[t]=ks.perm(X,Y,alpha=0.05,B=200)  
}  
mean(decision)
```

```
## [1] 0.07
```

Check the Type II error:

```
n = 100; m=100  
times = 100  
decision <- numeric(times)  
for (t in 1:times) {  
  X=rnorm(n)  
  Y=rnorm(m, mean = 0.2)  
  decision[t]=ks.perm(X,Y,alpha=0.05,B=200)  
}  
mean(1-decision)
```

```
## [1] 0.79
```

Here is the nearest neighbor test function that I got from the lecture.

```
library(RANN)  
NNT <- function(z, ix, sizes, R) {  
  n1 <- sizes[1]  
  n2 <- sizes[2]  
  n <- n1 + n2  
  z <- z[ix, ]
```

```

NN <- nn2(z, z, k=R+1)
block1 <- NN$nn.idx[1:n1, -1]
block2 <- NN$nn.idx[(n1+1):n, -1]
i1 <- sum(block1 < n1 + .5)
i2 <- sum(block2 > n1 + .5)
return((i1 + i2) / (R * n))
}

NNT.perm <- function(z, sizes, R,alpha,B) {
  T0=NNT(z,1:nrow(z),sizes, R)
  TPerm=rep(0,B)
  for (b in 1:B) {
    permindex=sample(1:nrow(z))
    TPerm[b]=NNT(z,permindex,sizes, R)
  }
  P <- mean(c(T0, TPerm) >= T0)
  return(P<alpha)
}

```

Check the Type I error:

```

n = 100; m=100; d=1
times = 100
decision <- numeric(times)
for (t in 1:times) {
  x <- matrix(rnorm(n*d), n, d)
  y <- matrix(rnorm(m*d), m, d)
  z <- rbind(x, y)
  decision[t]=NNT.perm(z, c(n,m), R=3,alpha=0.05,B=500)
}
mean(decision)

```

```
## [1] 0.09
```

Check the Type II error:

```

n = 100; m=100; d=1
times = 100
decision <- numeric(times)
for (t in 1:times) {
  x <- matrix(rnorm(n*d), n, d)
  y <- matrix(rnorm(m*d, mean = 0.4), m, d)
  z <- rbind(x, y)
  decision[t]=NNT.perm(z, c(n,m), R=3,alpha=0.05,B=200)
}
mean(1-decision)

```

```
## [1] 0.83
```

- **2b.** The choice of distance is important in distance correlation. Besides Euclidean distance, we can also consider  $L_p$  distance

$$\|X - Y\|_{\ell_p} = \left( \sum_{i=1}^d |X_i - Y_i|^p \right)^{1/p}.$$

Can you design a simulation experiment to assess the Type I and Type II error of distance correlation test for  $p = 1, 2, 4, 8$ .

```
plst = c(1,2,4,8)
```

Define the needed functions:

```
library(MASS)
Akl <- function(x, pow) {
  d <- as.matrix(dist(x, method="minkowski", p=pow))
  m <- rowMeans(d)
  M <- mean(d)
  a <- sweep(d, 1, m)
  b <- sweep(a, 2, m)
  return(b + M)
}
DCOR <- function(x, y, pow) {
  x <- as.matrix(x)
  y <- as.matrix(y)
  n <- nrow(x)
  m <- nrow(y)
  if (n != m || n < 2) stop("Sample sizes must agree")
  if (! (all(is.finite(c(x, y)))))
    stop("Data contains missing or infinite values")
  A <- Akl(x, pow)
  B <- Akl(y, pow)
  dCov <- sqrt(mean(A * B))
  dVarX <- sqrt(mean(A * A))
  dVarY <- sqrt(mean(B * B))
  dCor <- sqrt(dCov / sqrt(dVarX * dVarY))
  list(dCov=dCov, dCor=dCor, dVarX=dVarX, dVarY=dVarY)
}
DCOR.perm <- function(x,y,alpha,B, pow) {
  distx <- as.matrix(dist(x, method="minkowski", p=pow))
  disty <- as.matrix(dist(y, method="minkowski", p=pow))
  T0=DCOR(distx, disty, pow)$dCor
  TPerm=rep(0,B)
  for (b in 1:B) {
    permindex=sample(1:nrow(x))
    permdisty=disty[permindex,permindex]
    TPerm[b]=DCOR(distx, permdisty, pow)$dCor
  }
  P <- mean(c(T0, TPerm) >= T0)
  return(P<alpha)
}
```

We can do the Distance correlation test by using the above function, and print the result.

```
z <- as.matrix(iris[1:5, 1:4])
x <- z[, 1:2]
y <- z[, 3:4]
for(p in plst) {
  print(c("when p is", p))
```

```

# Get Type I
n = 100; mu=rep(0,6); sigma=diag(6)
times = 20
decision <- numeric(times)
for (t in 1:times) {
  data <- mvrnorm(n, mu, sigma)
  x <- data[,1:3]
  y <- data[,4:6]
  decision[t]=DCOR.perm(x, y,alpha=0.05,B=100, p)
}

print(c("Type I error is:", mean(decision)))

# Get Type II
rho=0.1
n = 100; mu=rep(0,6); sigma=diag(6)*(1-rho)+rho
times = 20
decision <- numeric(times)
for (t in 1:times) {
  data <- mvrnorm(n, mu, sigma)
  x <- data[,1:3]
  y <- data[,4:6]
  decision[t]=DCOR.perm(x, y,alpha=0.05,B=100, p)
}
print(c("Type II error is:", mean(1-decision)))
print("-----")
}

```

```

## [1] "when p is" "1"
## [1] "Type I error is:" "0"
## [1] "Type II error is:" "0.75"
## [1] "-----"
## [1] "when p is" "2"
## [1] "Type I error is:" "0.05"
## [1] "Type II error is:" "0.5"
## [1] "-----"
## [1] "when p is" "4"
## [1] "Type I error is:" "0.05"
## [1] "Type II error is:" "0.6"
## [1] "-----"
## [1] "when p is" "8"
## [1] "Type I error is:" "0.05"
## [1] "Type II error is:" "0.75"
## [1] "-----"

```

- **2c.** The Kendall tau's correlation can also be used to test the independence between variables when they are both one-dimensional variable. Kendall tau's correlation is defined as

$$\frac{1}{n(n-1)} \sum_{i \neq j} \text{sgn}(x_i - x_j) \text{sgn}(y_i - y_j).$$

You can use the implementation in `cor` by setting `method="kendall"`. Can you implement an permutation-based independence test for Kendall tau's correlation? Specifically, write a function

kendalltau.perm(x,y,alpha,B), where x and y are input data, alpha is significance level and B is number of replicate in permutation test. kendalltau.perm(x,y,alpha,B) returns the decision on whether the null hypothesis is rejected.

```
kendalltau.perm <- function(x,y,alpha,B) {
  T0=cor.test(x, y,method="kendall")$statistic
  TPerm=rep(0,B)
  for (b in 1:B) {
    permx=sample(1:length(y))
    permy=y[permx]
    TPerm[b]=cor.test(x, permy,method="kendall")$statistic
  }
  P <- mean(c(T0, TPerm) >= T0)
  return(P<alpha)
}
x = rnorm(100)
y = rnorm(100)
kendalltau.perm(x, y, 0.05, 100)
```

```
## [1] FALSE
```

- **2d.** Both distance correlation test and Kendall tau's correlation test can be used to test the independence between one-dimensional variables. Can you design a simulation experiment to compare Type I and Type II error of distance correlation test and Kendall tau's correlation test? Yes, I can.

```
n = 100; m=100
times = 60
decision.DCOR <- numeric(times)
decision.KEN <- numeric(times)
for (t in 1:times) {
  X=as.matrix(rnorm(n))
  Y=as.matrix(rnorm(m))
  decision.DCOR[t]=DCOR.perm(X,Y,alpha=0.05,B=100, 2)
  decision.KEN[t]=kendalltau.perm(X, Y,alpha=0.05,B=100)
}
print(c("Type I error of DCOR is:", mean(decision.DCOR)))
```

```
## [1] "Type I error of DCOR is:" "0.0833333333333333"
```

```
print(c("Type I error of kendalltau is:", mean(decision.KEN)))
```

```
## [1] "Type I error of kendalltau is:" "0.0833333333333333"
```

```
# Get Type II
```

```
n = 100; m=100
times = 100
decision.DCOR2 <- numeric(times)
decision.KEN2 <- numeric(times)
for (t in 1:times) {
  X=as.matrix(rnorm(n))
```

```
Y=as.matrix(rnorm(m))
decision.DCOR2[t]=DCOR.perm(X,Y,alpha=0.05,B=100, 2)
decision.KEN2[t]=kendalltau.perm(X,Y,alpha=0.05,B=100)
}
print(c("Type II error of DCOR is:", mean(1-decision.DCOR2)))
```

```
## [1] "Type II error of DCOR is:" "0.96"
```

```
print(c("Type II error of kendalltau is:", mean(1-decision.KEN2)))
```

```
## [1] "Type II error of kendalltau is:" "0.97"
```